

MARKOV CHAINS 12

COMPUTER SCIENCE MENTORS 70

November 26 - 30, 2018

1 Markov Chains

P is a **transition probability matrix** if:

1. All of the entries are non-negative.
2. The sum of entries in each row is 1.

A **Markov chain** is defined by four things: $(\mathcal{X}, \pi_0, P, \{X_n\}_{n=0}^\infty)$

\mathcal{X}	Set of states
π_0	Initial probability distribution (the chance you are in any of the states)
P	Transition probability matrix
$\{X_n\}_{n=0}^\infty$	Sequence of random variables where:
	$P[X_0 = i] = \pi_0(i), i \in \mathcal{X}$
	$P[X_{n+1} = j X_n = i, X_{n-1}, \dots, X_0] = P(i, j), \forall n \geq 0, \forall i, j \in \mathcal{X}$

A Markov chain is **irreducible** if we can go from any state to any other state, possibly in multiple steps.

Periodicity has to do with the period of occurrence of a state. If a state s has period 2, the Markov chain can be in s at every other time point. If a state has period 1, it's aperiodic; otherwise, it's periodic. More quantitatively, define value $d(i)$ for each state i as:

$$d(i) := \text{g.c.d}\{n > 0 | P^n(i, i) = P[X_n = i | X_0 = i] > 0\}, i \in \mathcal{X}$$

If $d(i) = 1$, then the Markov chain is **aperiodic**. If $d(i) \neq 1$, then the Markov chain is periodic and its **period** is $d(i)$.

A distribution π is **invariant** for the transition probability P if it satisfies the following **balance equations**


$$\pi \cdot P = \pi.$$

Theorem 1: A finite irreducible Markov chain has a unique invariant distribution.

Theorem 2: All irreducible and aperiodic Markov chains converge to the unique invariant distribution - i.e, after a large amount of steps, the chance you are in any given state is given by the invariant distribution.

Theorem 3: If a Markov chain is finite and reducible, the amount of time spent in each state approaches the invariant distribution as n grows large

Equations that model what will happen at the next step are called **first step equations**




Denote $\beta(i, j)$ as the expected amount of time it would take to move from i to j . $\beta(0, 1) = 1 + \frac{1}{2} \cdot \beta(0, 1)$ $\beta(1, 1) = 0$

1.1 Questions


1. Types of Markov Chains

- (a) Draw a Markov Chain that is reducible.



- (b) Draw a periodic Markov Chain (note that when we talk about periodicity, we imply that the Markov chain is reducible). Does it have an invariant distribution? Does it always converge to this invariant distribution? Does the invariant distribution represent the average amount of time spent in each state in this chain?


Solution: For the following example, we see that this chain is periodic because in order to get back to a state, you must travel some multiple of 2 steps. For these types of Markov chains, it **does not always converge** to the invariant distribution. Clearly, in this case the invariant distribution is $\pi = [0.5 \ 0.5]$. However, if we start with the distribution $\pi_0 = [1 \ 0]$, this distribution will never converge to π . The invariant distribution **does represent the average amount of time spent in each state**, because this chain is finite and



- (c) Draw an aperiodic Markov Chain. Does it have an invariant distribution? Does it converge to this invariant distribution? Does the invariant distribution represent

the average amount of time spent in each state in this chain?


Solution: For these types of Markov chains, it **always** to the invariant distribution. The invariant distribution **does represent the average amount of time spent in each state**, because this chain is finite and irreducible.



2. Life of Alex

Alex is enjoying college life. She spends a day either studying, partying, or looking for housing for the next year. If she is studying, the chances of her studying the next day are 30%, the chances of her partying the next day are 50%, and the chances of her looking for housing the next day are 20%. If she is partying, the chances of her partying the next day are 10%, the chances of her studying the next day are 60%, and the chances of her looking for housing the next day are 30%. If she is looking for housing, the chances of her looking for housing the next day are 50%, the chances of her partying the next day are 30% and the chances of her studying the next day are 20%.

(a) Draw a Markov chain to visualize Alex's life.



Solution:

(b) Write out a matrix to represent this Markov chain.

Solution:

$$\begin{bmatrix} .3 & .5 & .2 \\ .6 & .1 & .3 \\ .2 & .3 & .5 \end{bmatrix}$$

- (c) If Alex studies on Monday, what is the chance that she is partying on Friday? (Don't do the math, just write out the expression that you would use to find it.)

Solution: If P is the matrix above, then it is $[1, 0, 0] \cdot P^4$

- (d) What percentage of her time should Alex expect to use looking for housing?

Solution: Solve the following system of equations: (first step equations)

$$S = .3S + .6P + .2H$$

$$P = .5S + .1P + .3H$$

$$H = .2S + .3P + .5H$$

$$S + P + H = 1$$

- (e) If Alex parties on Monday, what is the chance of Alex partying again before studying?

Solution: Set up the following equations:

$$H1 = 0$$

$$H2 = .6(H1) + .1(1) + .3(H3)$$

$$H3 = .2(H1) + .3(1) + .5(H3)$$

Solving for $H2$, we get 0.28.

3. Stanford Cinema

You have a database of an infinite number of movies. Each movie has a rating that is uniformly distributed in 0, 1, 2, 3, 4, 5 independent of all other movies. You want to find two movies such that the sum of their ratings is greater than 7.5 (7.5 is not included).

- a) A Stanford student chooses two movies each time and calculates the sum of their ratings. If it is less than or equal to 7.5, the student throws away these two movies

and chooses two other movies. The student stops when he/she finds two movies such that the sum of their ratings is greater than 7.5. What is the expected number of movies that this student needs to choose from the database?

Solution: Each time when the Stanford student chooses two movies, there are $6^2 = 36$ different possible pairs of ratings. We know that there are 4 pairs whose sum is greater than 7.5. Therefore, the probability that in a single trial, the Stanford student gets two movies such that the sum of their ratings is greater than 7.5 is $\frac{1}{9}$. Then the number of times that the student needs to pick movies is geometrically distributed with mean 9. Then the expected number of movies that the student needs to choose is 18.

- b) A Berkeley student chooses movies from the database one by one and keeps the movie with the highest rating. The student stops when he/she finds the sum of the ratings of the last movie that he/she has chosen and the movie with the highest rating among all the previous movies is greater than 7.5. What is the expected number of movies that the student will have to choose?

Solution: To solve this problem, let's create a Markov chain with 5 states. One state, S_0 will represent starting or having a movie with rating 0, 1 or 2 in hand. This state is special because no matter what the next state is, we won't be able to draw a movie on the next round that will have the sum of the ratings be more than 8. We will have a "Done" state, S_4 that represents finding 2 movies that have rating more than 8. The other 3 states S_1, S_2, S_3 will represent having movies with rating 3, 4, 5 in hand, respectively. From each of these states, note that we have some probability of drawing some movie on the next step that puts the total greater than 8. This is the transition matrix representing the chain:

$$\begin{bmatrix} 1/2 & 1/6 & 1/6 & 1/6 & 0 \\ 0 & 2/3 & 1/6 & 0 & 1/6 \\ 0 & 0 & 2/3 & 0 & 1/3 \\ 0 & 0 & 0 & 1/2 & 1/2 \\ 0 & 0 & 0 & 0 & 1 \end{bmatrix}$$

The corresponding hitting time equations are:

$$\begin{aligned} \beta(4) &= 0 \\ \beta(3) &= \frac{1}{2}\beta(3) + \frac{1}{2}\beta(4) \\ \beta(2) &= \frac{2}{3}\beta(3) + \frac{1}{3}\beta(4) \end{aligned}$$

$$\beta(1) = \frac{2}{3}\beta(1) + \frac{1}{6}\beta(3) + \frac{1}{6}\beta(3)$$

$$\beta(0) = \frac{1}{2}\beta(0) + \frac{1}{6}\beta(1) + \frac{1}{6}\beta(2) + \frac{1}{6}\beta(3)$$


Solving these, we find $\beta(0) = \frac{31}{6}$. This shows that the Berkeley student is smarter than the Stanford student.

4. Bet On It

Smith is in jail and has 3 dollars; he can get out on bail if he has 8 dollars. A guard agrees to make a series of bets with him. If Smith bets A dollars, he wins A dollars with probability 0.4 and loses A dollars with probability 0.6.

- a) Find the probability that he wins 8 dollars before losing all of his money if he bets 1 dollar each time.

Solution: The Markov chain $(X_n, n = 0, 1, \dots)$ representing the evolution of Smith's money has diagram



Let $\phi(i)$ be the probability that the chain reaches state 8 before reaching state 0, starting from state i . In other words, if S_j is the first $n \geq 0$ such that $X_n = j$,

$$P_i(S_8 < S_0) = P(S_8 < S_0 | X_0 = i) \quad (1)$$

Using first-step analysis (viz. the Markov property at time $n = 1$), we have

$$\phi(i) = 0.4\phi(i+1) + 0.6\phi(i-1), \quad i = 1, 2, 3, 4, 5, 6, 7$$

$$\phi(0) = 0$$

$$\phi(8) = 1.$$


We solve this system of linear equations and find

$$\begin{aligned} \phi &= (\phi(1), \phi(2), \phi(3), \phi(4), \phi(5), \phi(6), \phi(7)) \\ &= (0.0203, 0.0508, 0.0964, 0.1649, 0.2677, 0.4219, 0.6531, 1). \end{aligned}$$

E.g. the probability that the chain reaches state 8 before reaching state 0, starting from state 3, is the third component of the vector and is equal to 0.0964. Note that $\phi(i)$ is increasing in i , which was expected.

- b) Find the probability that he wins 8 dollars before losing all of his money if he bets, each time, as much as possible but not more than necessary to bring his fortune up to 8 dollars.

Solution: Now the chain is:



and the equations are:

$$\phi(3) = 0.4\phi(6)$$

$$\phi(6) = 0.4\phi(8) + 0.6\phi(4)$$

$$\phi(4) = 0.4\phi(8)$$

$$\phi(0) = 0$$

$$\phi(8) = 1$$

We solve and find

$$\phi(3) = 0.256, \phi(4) = 0.4, \phi(6) = 0.64$$

- c) Which strategy gives Smith the better chance of getting out of jail?


Solution: By comparing the third components of the vector ϕ we find that the bold strategy gives Smith a better chance to get out of jail.

5. Tossing Coins

A fair coin is tossed repeatedly and independently. Find the expected number of tosses till the pattern HTH appears.

Solution: Call HTH our target. Consider a chain that starts from a state called nothing (\emptyset) and is eventually absorbed at HTH. If we first toss H then we move to state H because this is the first letter of our target. If we toss a T then we move back to \emptyset having expended 1 unit of time. Being in state H we either move to a

new state HT if we bring T and we are 1 step closer to the target or, if we bring H, we move back to H: we have expended 1 unit of time, but the new H can be the beginning of a target. When in state HT we either move to HTH and we are done or, if T occurs then we move to \emptyset . The transition diagram is



Let $\phi(i)$ be the expected number of steps to reach HTH starting from i . We have

$$\phi(HT) = 1 + \frac{1}{2}\phi(\emptyset)$$

$$\phi(H) = 1 + \frac{1}{2}\phi(H) + \frac{1}{2}\phi(HT)$$

$$\phi(\emptyset) = 1 + \frac{1}{2}\phi(\emptyset) + \frac{1}{2}\phi(H)$$

We solve and find $\phi(\emptyset) = 10$.